

Predicting On Base Percentage (OBP)

W. Kerney

2024

Project Overview

This project focuses on designing, building, and evaluating a model to predict a player's on-base percentage (OBP) based on other performance indicators. The MLB Hitting and Pitching stats dataset contains individual statistics for Major League Baseball players. It includes various performance metrics, such as batting averages, on-base percentages, slugging percentages, home runs, RBIs, strikeouts, walks, and more.

Step 1: Load and Explore the Dataset

```
install.packages("readr")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```
install.packages("tidyverse")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```
install.packages("dplyr")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```
install.packages("janitor")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```
install.packages("caret")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```
install.packages("randomForest")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```
install.packages("knitr")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```
install.packages("Metrics")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.4'  
## (as 'lib' is unspecified)
```

```

library(readr)
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v purrr      1.0.2
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(janitor)

##
## Attaching package: 'janitor'
##
## The following objects are masked from 'package:stats':
##
##   chisq.test, fisher.test

library(knitr)
library(Metrics)

#read in data
mlb_data <- read.csv("baseball_hitting.csv") %>%
  clean_names()

# Explore the dataset
head(mlb_data)

##   player_name position games at_bat runs hits double_2b third_baseman home_run
## 1      B Bonds      LF  2986   9847 2227 2935         601           77       762
## 2      H Aaron      RF  3298  12364 2174 3771         624           98       755
## 3      B Ruth      RF  2504   8399 2174 2873         506          136       714
## 4      A Pujols     1B  3080  11421 1914 3384         686           16       703
## 5 A Rodriguez     SS  2784  10566 2021 3115         548           31       696
## 6      W Mays      CF  2992  10881 2062 3283         523          140       660
##   run_batted_in a_walk strikeouts stolen_base caught_stealing   avg
## 1          1996   2558         1539         514          141 0.298
## 2          2297   1402         1383         240           73 0.305
## 3          2213   2062         1330         123          117 0.342
## 4          2218   1373         1404         117           43 0.296
## 5          2086   1338         2287         329           76 0.295
## 6          1903   1464         1526         338          103 0.302
##   on_base_percentage slugging_percentage on_base_plus_slugging
## 1              0.444              0.607              1.051
## 2              0.374              0.555              0.929
## 3              0.474              0.690              1.164
## 4              0.374              0.544              0.918
## 5              0.380              0.550              0.930
## 6              0.384              0.557              0.941

```

```
summary(mlb_data)
```

```
## player_name      position      games      at_bat
## Length:2508      Length:2508      Min.   : 2.0      Min.   : 262
## Class :character  Class :character  1st Qu.: 616.8    1st Qu.: 1874
## Mode  :character  Mode  :character  Median : 998.0    Median : 3266
##                                     Mean  :1084.6      Mean   : 3715
##                                     3rd Qu.:1438.2    3rd Qu.: 5106
##                                     Max.   :3562.0    Max.   :14053
##                                     NA's   :8          NA's   :8
##      runs      hits      double_2b      third_base
## Min.   : 32.0    Min.   : 57.0    Min.   : 7.0     Min.   : 0.00
## 1st Qu.: 231.8    1st Qu.: 471.2    1st Qu.: 86.0    1st Qu.: 9.00
## Median : 423.5    Median : 853.5    Median :154.0    Median : 20.00
## Mean   : 521.6    Mean   :1010.9    Mean   :181.9    Mean   : 32.33
## 3rd Qu.: 719.2    3rd Qu.:1399.2    3rd Qu.:249.0    3rd Qu.: 42.25
## Max.   :2295.0    Max.   :4256.0    Max.   :792.0    Max.   :309.00
## NA's   :8        NA's   :8        NA's   :8        NA's   :8
##      home_run      run_batted_in      a_walk      strikeouts
## Min.   : 17.0      Min.   : 37.0      Min.   : 19.0      Length:2508
## 1st Qu.: 33.0      1st Qu.: 222.0     1st Qu.: 162.0     Class :character
## Median : 69.0      Median : 404.0     Median : 292.5     Mode  :character
## Mean   :100.6      Mean   : 494.2     Mean   : 373.0
## 3rd Qu.:125.5      3rd Qu.: 656.2     3rd Qu.: 486.2
## Max.   :762.0      Max.   :2297.0     Max.   :2558.0
## NA's   :8        NA's   :8        NA's   :8
##      stolen_base      caught_stealing      avg      on_base_percentage
## Min.   : 0.0          Length:2508      Min.   :0.1230      Min.   :0.1570
## 1st Qu.: 11.0          Class :character  1st Qu.:0.2470      1st Qu.:0.3110
## Median : 32.0          Mode  :character  Median :0.2620      Median :0.3300
## Mean   : 76.1                                     Mean  :0.2633      Mean   :0.3316
## 3rd Qu.: 89.0                                     3rd Qu.:0.2780      3rd Qu.:0.3510
## Max.   :1406.0                                     Max.   :0.3670      Max.   :0.4820
## NA's   :8          NA's   :8        NA's   :8
##      slugging_percentage      on_base_plus_slugging
## Min.   :0.1970      Min.   :0.3540
## 1st Qu.:0.3750      1st Qu.:0.6930
## Median :0.4070      Median :0.7380
## Mean   :0.4099      Mean   :0.7417
## 3rd Qu.:0.4410      3rd Qu.:0.7840
## Max.   :0.6900      Max.   :1.1640
## NA's   :8          NA's   :20
```

Step 2: Data Preprocessing and Feature Engineering

Now we'll extract relevant features and create new ones to use in the model.

```
#some feature engineering
library(dplyr)

#converting strikeouts to integers
mlb_data$strikeouts <- as.integer(mlb_data$strikeouts)
```

```
## Warning: NAs introduced by coercion
```

```
str(mlb_data)
```

```
## 'data.frame':    2508 obs. of  18 variables:
##  $ player_name      : chr  "B Bonds" "H Aaron" "B Ruth" "A Pujols" ...
##  $ position         : chr  "LF" "RF" "RF" "1B" ...
##  $ games            : int   2986 3298 2504 3080 2784 2992 2671 2543 2354 2808 ...
##  $ at_bat            : int   9847 12364 8399 11421 10566 10881 9801 8422 8813 10006 ...
##  $ runs              : int   2227 2174 2174 1914 2021 2062 1662 1583 1475 1829 ...
##  $ hits              : int   2935 3771 2873 3384 3115 3283 2781 2328 2408 2943 ...
##  $ double_2b         : int    601 624 506 686 548 523 524 451 379 528 ...
##  $ third_baseaman    : int    77 98 136 16 31 140 38 26 45 72 ...
##  $ home_run          : int    762 755 714 703 696 660 630 612 609 586 ...
##  $ run_batted_in     : int   1996 2297 2213 2218 2086 1903 1836 1699 1667 1812 ...
##  $ a_walk            : int   2558 1402 2062 1373 1338 1464 1312 1747 929 1420 ...
##  $ strikeouts         : int   1539 1383 1330 1404 2287 1526 1779 2548 2306 1532 ...
##  $ stolen_base       : int    514 240 123 117 329 338 184 19 234 204 ...
##  $ caught_stealing   : chr   "141" "73" "117" "43" ...
##  $ avg               : num   0.298 0.305 0.342 0.296 0.295 0.302 0.284 0.276 0.273 0.294 ...
##  $ on_base_percentage : num   0.444 0.374 0.474 0.374 0.38 0.384 0.37 0.402 0.344 0.389 ...
##  $ slugging_percentage : num   0.607 0.555 0.69 0.544 0.55 0.557 0.538 0.554 0.534 0.537 ...
##  $ on_base_plus_slugging: num   1.051 0.929 1.164 0.918 0.93 ...
```

```
#creating new features
```

```
mlb_features <- mlb_data %>%
```

```
  filter(at_bat > 0) %>%
```

```
  mutate(ISO = slugging_percentage - ((hits/ at_bat)),
```

```
         BABIP = (hits - home_run) / (at_bat - strikeouts - home_run)) %>%
```

```
select(player_name, position, on_base_percentage, ISO, BABIP, slugging_percentage, runs, run_batted_in,
```

```
# Remove rows with NA values and convert column names.
```

```
mlb_features <- mlb_features[complete.cases(mlb_features), ]
```

```
colnames(mlb_features)[c (3,6,7,8,9,10,11,12)] <- c("OBP", "SLG", "R", "RBI", "K", "BB", "HR", "OPS")
```

```
# Inspect the new dataset
```

```
head(mlb_features)
```

```
##  player_name position  OBP      ISO      BABIP  SLG    R  RBI    K  BB  HR
##  1      B Bonds      LF 0.444 0.3089397 0.2879671 0.607 2227 1996 1539 2558 762
##  2      H Aaron      RF 0.374 0.2500016 0.2949345 0.555 2174 2297 1383 1402 755
##  3      B Ruth      RF 0.474 0.3479355 0.3397325 0.690 2174 2213 1330 2062 714
##  4      A Pujols     1B 0.374 0.2477037 0.2878463 0.544 1914 2218 1404 1373 703
##  5 A Rodriguez     SS 0.380 0.2551864 0.3190030 0.550 2021 2086 2287 1338 696
##  6      W Mays      CF 0.384 0.2552814 0.3016676 0.557 2062 1903 1526 1464 660
##      OPS
##  1 1.051
##  2 0.929
##  3 1.164
##  4 0.918
##  5 0.930
##  6 0.941
```

Step 3: Splitting the Data

Now that we have confirmed that the features we created are valid. We'll split the data, creating a training set and a test set for evaluating the model.

```
# Split the data into training and test sets
library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'
## The following objects are masked from 'package:Metrics':
##
##   precision, recall
## The following object is masked from 'package:purrr':
##
##   lift

set.seed(123)
trainIndex <- createDataPartition(mlb_features$OBP, p = .8, list = FALSE)

mlb_train <- mlb_features[trainIndex, ]
mlb_test  <- mlb_features[-trainIndex, ]
```

Step 4: Train a Machine Learning Model

After splitting our data, we'll use a random forest model to predict OBP.

```
#creating a random forest model to predict OBP
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:dplyr':
##
##   combine
## The following object is masked from 'package:ggplot2':
##
##   margin

# Train the random forest model
rf_model <- randomForest(OBP ~ ISO + BABIP + SLG + R + RBI + K + BB + HR, data = mlb_train, ntree = 100)

# Check the model's summary
print(rf_model)


##
## Call:
## randomForest(formula = OBP ~ ISO + BABIP + SLG + R + RBI + K + BB + HR, data = mlb_train, ntree = 100)
##           Type of random forest: regression
##           Number of trees: 100
## No. of variables tried at each split: 2
```

```
##  
##           Mean of squared residuals: 0.0001547024  
##           % Var explained: 83.93
```

Step 5: Evaluate the Model's Performance

From here we'll assess the model's performance using the test data and visualize the results.

```
#evaluating model performance  
# Make predictions on the test set  
obp_predictions <- predict(rf_model, mlb_test)  
rmse(predict(rf_model, mlb_test), mlb_test$OBP)  
  
## [1] 0.01065553  
  
# Calculate the R-squared value  
rsq <- cor(mlb_test$OBP, obp_predictions)^2  
print(paste("R-squared:", round(rsq, 2)))  
  
## [1] "R-squared: 0.88"  
  
# Visualize predicted vs. actual OBP  
library(ggplot2)  
  
ggplot(mlb_test, aes(x = OBP, y = obp_predictions)) +  
  geom_point() +  
  geom_abline(slope = 1, intercept = 0, linetype = "dashed", color = "red") +  
  labs(  
    title = "Actual vs. Predicted OBP",  
    x = "Actual OBP",  
    y = "Predicted OBP"  
  )
```



Tuning Model

The Mean of squared residuals and % Var explained lets us know we do not have an overfit model, and evaluating the rmse we have our base error so now we will use caret to do some tuning.

```
# use caret to pick a value for mtry

tuned_model <- train(OBP ~ ISO + BABIP + SLG + R + RBI + K + BB + HR ,data = mlb_train, ntree = 50, #nu
                     method = "rf")


print(tuned_model)

## Random Forest
##
## 1992 samples
##    8 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1992, 1992, 1992, 1992, 1992, 1992, ...
## Resampling results across tuning parameters:
##
##  mtry  RMSE      Rsquared  MAE
##    2   0.01315634 0.8484628 0.009724142
##    5   0.01147130 0.8792759 0.008347823
##    8   0.01138418 0.8757369 0.008278651
##
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final value used for the model was mtry = 8.
```

As we can see, caret picked 8 as the best mtry value, letting us know that the model was most accurate when we use 8 variables for prediction. We can see this very clearly if we plot the rmse for each of the values it tried by passing our tuned model right to ggplot.

```
# plot the rmse for various possible training values  
ggplot(tuned_model)
```



Comparing the Models

From here we will check the model against the test data.

```
#Checking original against tuned model  
print("base model rmse:")
```

```
## [1] "base model rmse:"
```

```
print(rmse(predict(rf_model, mlb_test), mlb_test$OBP))
```

```
## [1] 0.01065553
```

```
print("tuned model rmse:")
```

```
## [1] "tuned model rmse:"
```

```
print(rmse(predict(tuned_model$finalModel, mlb_test), mlb_test$OBP))
```

```
## [1] 0.009343885
```

We can also compare what each model selected as the most important features. In this case, we are going to look at the top five features for each model. In the plots, 'IncNodePurity' is a measure of how important

each feature is. A larger value means that feature was more important.

```
# plot both plots at once
par(mfrow = c(1,2))

varImpPlot(rf_model, n.var = 5)
varImpPlot(tuned_model$finalModel, n.var = 5)
```

